

Data Science Language Analysis

Statistical Analysis

Nazli Ozum Kafae

Prash Medirattaa

Avinash Prabhakaran

2018-04-15

```
#Loading the required packages
suppressPackageStartupMessages(library(tidyverse))
suppressPackageStartupMessages(library(knitr))

#Reading in processed data.
responses <- read.csv(file = "../docs/survey_results_clean.csv")

#preference encoding the response variable. Python -> 1; R -> 0
data <- responses %>% mutate(preference = if_else(preference == "Python", 1, 0))

data$background <- as.character(data$background)
data <- data %>%
  mutate(background = ifelse(background == "Computer Science / Computer Engineering",
                             "Computer Sc/Eng",
                             ifelse(background == "Mathematics / Statistics",
                                     "Maths/Stats", background)))

#Releveling the reference task from Data Viz -> Machine Learning
data_relevel <- data
data_relevel$task <- relevel(data$task, ref="Machine Learning")

#Fitting a GLM without any confounding variables.
mod <- glm(preference ~ task, family = binomial(link = 'logit'), data = data)
summary(mod)

##
## Call:
## glm(formula = preference ~ task, family = binomial(link = "logit"),
##      data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9623  -0.8519   0.5615   0.5615   1.5425
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -0.8267     0.4532  -1.824   0.0681 .
## taskData wrangling     0.8267     0.7008   1.180   0.2381
## taskMachine Learning  2.5943     0.6104   4.250 2.14e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
```

```

##      Null deviance: 110.372  on 84  degrees of freedom
## Residual deviance:  87.555  on 82  degrees of freedom
## AIC: 93.555
##
## Number of Fisher Scoring iterations: 4
#Fitting GLM with all the confounding variables.
mod <- glm(preference ~ task + background + experience + attitude + first + active,
            family = binomial(link = 'logit'), data = data)
summary(mod)

##
## Call:
## glm(formula = preference ~ task + background + experience + attitude +
##      first + active, family = binomial(link = "logit"), data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.90837  -0.37589   0.03797   0.36210   2.55585
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -5.3267     3.7988  -1.402   0.1609
## taskData wrangling      2.1788     1.2077   1.804   0.0712 .
## taskMachine Learning    3.5525     1.2092   2.938   0.0033 **
## backgroundComputer Sc/Eng 3.6830     1.7474   2.108   0.0351 *
## backgroundEngineering  -0.3603     1.5457  -0.233   0.8157
## backgroundMaths/Stats    0.6773     1.2335   0.549   0.5829
## backgroundOther          0.5810     1.4710   0.395   0.6929
## experienceLess than 1   -0.5551     0.9276  -0.598   0.5495
## experienceMore than 5   -3.0570     2.1747  -1.406   0.1598
## attitudeNo             -0.9623     2.1569  -0.446   0.6555
## attitudeYes             1.4492     1.2635   1.147   0.2514
## firstJava              -2.0377     1.6539  -1.232   0.2179
## firstMatlab             0.3319     1.6907   0.196   0.8444
## firstOther              0.8743     1.3958   0.626   0.5311
## firstPython             4.0073     2.0407   1.964   0.0496 *
## firstR                 -0.9978     1.5098  -0.661   0.5087
## firstSAS                -0.7765     1.9379  -0.401   0.6886
## active2                  1.0669     2.7495   0.388   0.6980
## active3                  3.8766     3.0529   1.270   0.2042
## active4                  1.7934     3.1708   0.566   0.5717
## active5 or more         19.6718    2192.1364   0.009   0.9928
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 110.372  on 84  degrees of freedom
## Residual deviance:  48.394  on 64  degrees of freedom
## AIC: 90.394
##
## Number of Fisher Scoring iterations: 17

```

#Removing Attitude as Confounder as change

```
mod <- glm(preference ~ task + background + experience + first + active,
           family = binomial(link = 'logit'), data = data)
summary(mod)
```

```
##
## Call:
## glm(formula = preference ~ task + background + experience + first +
##      active, family = binomial(link = "logit"), data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6732  -0.3554   0.0497   0.4396   2.6319
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -4.5485     3.1287  -1.454  0.146006
## taskData wrangling      2.2772     1.1859   1.920  0.054824 .
## taskMachine Learning    4.1062     1.1574   3.548  0.000389 ***
## backgroundComputer Sc/Eng  3.9260     1.7344   2.264  0.023597 *
## backgroundEngineering  -0.7602     1.4649  -0.519  0.603812
## backgroundMaths/Stats    0.2319     1.1111   0.209  0.834691
## backgroundOther         0.7808     1.3918   0.561  0.574814
## experienceLess than 1    -0.7904     0.9347  -0.846  0.397758
## experienceMore than 5   -2.7743     2.0524  -1.352  0.176476
## firstJava              -2.3143     1.5965  -1.450  0.147179
## firstMatlab            -0.3195     1.6155  -0.198  0.843204
## firstOther             0.9801     1.3456   0.728  0.466374
## firstPython            3.2122     1.6282   1.973  0.048510 *
## firstR                -1.0210     1.4712  -0.694  0.487702
## firstSAS              -0.3420     1.8965  -0.180  0.856889
## active2                1.3570     2.1491   0.631  0.527737
## active3                4.0532     2.4874   1.629  0.103209
## active4                2.4884     2.7204   0.915  0.360346
## active5 or more       19.6832    2173.2657   0.009  0.992774
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 110.37  on 84  degrees of freedom
## Residual deviance:  50.42  on 66  degrees of freedom
## AIC: 88.42
##
## Number of Fisher Scoring iterations: 17
```

#Removing Experience as Confounder

```
mod <- glm(preference ~ task + background + first + active,
           family = binomial(link = 'logit'), data = data)
summary(mod)
```

```
##
## Call:
## glm(formula = preference ~ task + background + first + active,
##      family = binomial(link = "logit"), data = data)
```

```
##
## Deviance Residuals:
##      Min        1Q      Median        3Q        Max
## -1.71388  -0.41985   0.05763   0.36812   2.86063
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -4.9757     2.9379  -1.694 0.090333 .
## taskData wrangling      2.1236     1.1197   1.897 0.057872 .
## taskMachine Learning    4.1142     1.1776   3.494 0.000476 ***
## backgroundComputer Sc/Eng  3.8721     1.6800   2.305 0.021178 *
## backgroundEngineering  -0.3729     1.2800  -0.291 0.770819
## backgroundMaths/Stats    0.4743     1.0138   0.468 0.639855
## backgroundOther         0.4334     1.2456   0.348 0.727903
## firstJava             -2.0701     1.4490  -1.429 0.153113
## firstMatlab           -0.6067     1.5278  -0.397 0.691302
## firstOther            0.6676     1.2053   0.554 0.579665
## firstPython           3.0143     1.6049   1.878 0.060356 .
## firstR               -0.9333     1.4231  -0.656 0.511941
## firstSAS             -0.6028     1.7588  -0.343 0.731806
## active2              1.4009     2.1620   0.648 0.516992
## active3              3.8922     2.3830   1.633 0.102411
## active4              2.7216     2.6835   1.014 0.310493
## active5 or more      20.2089    2047.0820   0.010 0.992123
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 110.372  on 84  degrees of freedom
## Residual deviance:  52.694  on 68  degrees of freedom
## AIC: 86.694
##
## Number of Fisher Scoring iterations: 17
##
#Removing active as Confounder
mod <- glm(preference ~ task + background + first,
           family = binomial(link = 'logit'), data = data)
summary(mod)

##
## Call:
## glm(formula = preference ~ task + background + first, family = binomial(link = "logit"),
##      data = data)
##
## Deviance Residuals:
##      Min        1Q      Median        3Q        Max
## -2.1334  -0.5820   0.1590   0.5909   2.5368
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -1.97466     1.45591  -1.356 0.175003
## taskData wrangling      1.53852     0.96548   1.594 0.111042
## taskMachine Learning    3.23967     0.88929   3.643 0.000269 ***
## backgroundComputer Sc/Eng  3.09956     1.45163   2.135 0.032742 *
```

```
## backgroundEngineering      0.01428    1.20813    0.012 0.990568
## backgroundMaths/Stats      0.32001    0.96980    0.330 0.741417
## backgroundOther            0.58445    1.09830    0.532 0.594631
## firstJava                  -1.73120    1.25914   -1.375 0.169161
## firstMatlab                -0.29960    1.23445   -0.243 0.808236
## firstOther                 0.90223    1.05911    0.852 0.394286
## firstPython                1.50843    1.21850    1.238 0.215737
## firstR                    -1.50020    1.30055   -1.154 0.248703
## firstSAS                   -1.20225    1.66039   -0.724 0.469018
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 110.372  on 84  degrees of freedom
## Residual deviance:  64.428  on 72  degrees of freedom
## AIC: 90.428
##
## Number of Fisher Scoring iterations: 6
```

Not Removing active as the AIC score of the model increases from 86 to 90.

#Removing first as Confounder

```
mod <- glm(preference ~ task + background + active,
           family = binomial(link = 'logit'), data = data)
summary(mod)
```

```
##
## Call:
## glm(formula = preference ~ task + background + active, family = binomial(link = "logit"),
##      data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.8748  -0.4723   0.2731   0.7162   2.1546
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -1.5684     1.8226  -0.861   0.3895
## taskData wrangling     1.9769     0.9406   2.102   0.0356 *
## taskMachine Learning    3.3671     0.8528   3.948 7.88e-05 ***
## backgroundComputer Sc/Eng  2.6873     1.1777   2.282   0.0225 *
## backgroundEngineering   -0.2201     0.9315  -0.236   0.8132
## backgroundMaths/Stats    0.6468     0.8725   0.741   0.4585
## backgroundOther         0.5662     0.9426   0.601   0.5481
## active2              -1.2156     1.7419  -0.698   0.4853
## active3               0.4494     1.7776   0.253   0.8004
## active4              -0.4845     2.0996  -0.231   0.8175
## active5 or more       17.4542    2141.6796   0.008   0.9935
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 110.372  on 84  degrees of freedom
```

```
## Residual deviance: 68.411 on 74 degrees of freedom
## AIC: 90.411
##
## Number of Fisher Scoring iterations: 17
```

Not Removing first language as the AIC score of the model increases from 86 to 90.

```
#Removing background as Confounder
```

```
mod <- glm(preference ~ task + first + active,
           family = binomial(link = 'logit'), data = data)
summary(mod)
```

```
##
## Call:
## glm(formula = preference ~ task + first + active, family = binomial(link = "logit"),
##      data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9176  -0.5017   0.2333   0.6315   2.6352
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -1.3523     1.6305  -0.829  0.406906
## taskData wrangling    1.1946     0.9287   1.286  0.198362
## taskMachine Learning  3.1200     0.8826   3.535  0.000408 ***
## firstJava          -1.5078     1.1569  -1.303  0.192484
## firstMatlab        -0.8223     1.3228  -0.622  0.534186
## firstOther         -0.2259     1.0561  -0.214  0.830614
## firstPython         1.7076     1.1688   1.461  0.144007
## firstR             -1.8317     1.1531  -1.588  0.112182
## firstSAS           -1.5946     1.4252  -1.119  0.263195
## active2            -0.2567     1.4501  -0.177  0.859496
## active3             1.8231     1.6670   1.094  0.274128
## active4             1.3506     1.9144   0.706  0.480483
## active5 or more    18.2516    2317.0584   0.008  0.993715
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 110.372 on 84 degrees of freedom
## Residual deviance: 63.657 on 72 degrees of freedom
## AIC: 89.657
##
## Number of Fisher Scoring iterations: 17
```

Not Removing background as the AIC score of the model increases from 86 to 89.

```
#Model with first language, background and active
```

```
mod <- glm(preference ~ task + background + first + active,
           family = binomial(link = 'logit'), data = data)
summary(mod)
```

```
##
## Call:
## glm(formula = preference ~ task + background + first + active,
```

```

##      family = binomial(link = "logit"), data = data)
##
## Deviance Residuals:
##      Min        1Q      Median        3Q        Max
## -1.71388  -0.41985   0.05763   0.36812   2.86063
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -4.9757     2.9379  -1.694 0.090333 .
## taskData wrangling      2.1236     1.1197   1.897 0.057872 .
## taskMachine Learning    4.1142     1.1776   3.494 0.000476 ***
## backgroundComputer Sc/Eng  3.8721     1.6800   2.305 0.021178 *
## backgroundEngineering  -0.3729     1.2800  -0.291 0.770819
## backgroundMaths/Stats    0.4743     1.0138   0.468 0.639855
## backgroundOther         0.4334     1.2456   0.348 0.727903
## firstJava             -2.0701     1.4490  -1.429 0.153113
## firstMatlab           -0.6067     1.5278  -0.397 0.691302
## firstOther            0.6676     1.2053   0.554 0.579665
## firstPython           3.0143     1.6049   1.878 0.060356 .
## firstR               -0.9333     1.4231  -0.656 0.511941
## firstSAS             -0.6028     1.7588  -0.343 0.731806
## active2              1.4009     2.1620   0.648 0.516992
## active3              3.8922     2.3830   1.633 0.102411
## active4              2.7216     2.6835   1.014 0.310493
## active5 or more      20.2089    2047.0820   0.010 0.992123
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 110.372  on 84  degrees of freedom
## Residual deviance:  52.694  on 68  degrees of freedom
## AIC: 86.694
##
## Number of Fisher Scoring iterations: 17
#Releveled model with first language, background and active
model <- glm(preference ~ task + background + first + active,
             family = binomial(link = 'logit'), data = data_relevel)
summary(model)

##
## Call:
## glm(formula = preference ~ task + background + first + active,
##      family = binomial(link = "logit"), data = data_relevel)
##
## Deviance Residuals:
##      Min        1Q      Median        3Q        Max
## -1.71388  -0.41985   0.05763   0.36812   2.86063
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -0.8615     2.6528  -0.325 0.745369
## taskData visualization -4.1142     1.1776  -3.494 0.000476 ***
## taskData wrangling    -1.9906     0.9898  -2.011 0.044324 *

```

```

## backgroundComputer Sc/Eng      3.8721      1.6800      2.305 0.021178 *
## backgroundEngineering      -0.3729      1.2800     -0.291 0.770819
## backgroundMaths/Stats       0.4743      1.0138      0.468 0.639855
## backgroundOther            0.4334      1.2456      0.348 0.727903
## firstJava                  -2.0701      1.4490     -1.429 0.153113
## firstMatlab                -0.6067      1.5278     -0.397 0.691302
## firstOther                 0.6676      1.2053      0.554 0.579665
## firstPython                3.0143      1.6049      1.878 0.060356 .
## firstR                    -0.9333      1.4231     -0.656 0.511941
## firstSAS                  -0.6028      1.7588     -0.343 0.731806
## active2                    1.4009      2.1620      0.648 0.516992
## active3                    3.8922      2.3830      1.633 0.102411
## active4                    2.7216      2.6835      1.014 0.310493
## active5 or more            20.2089     2047.0820      0.010 0.992123
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 110.372  on 84  degrees of freedom
## Residual deviance:  52.694  on 68  degrees of freedom
## AIC: 86.694
##
## Number of Fisher Scoring iterations: 17
# Comparison of AIC scores for all the models
bind_cols(model = c("preference ~ task",
                    "preference ~ task + background + experience + attitude + first + active",
                    "preference ~ task + background + experience + first + active",
                    "preference ~ task + background + first + active",
                    "preference ~ task + background + first",
                    "preference ~ task + background + active",
                    "preference ~ task + first + active"),
          AIC = c(93.555, 90.394, 88.42, 86.694, 90.428, 90.411, 89.657)) %>%
kable()

```

| model | AIC |
|---|--------|
| preference ~ task | 93.555 |
| preference ~ task + background + experience + attitude + first + active | 90.394 |
| preference ~ task + background + experience + first + active | 88.420 |
| preference ~ task + background + first + active | 86.694 |
| preference ~ task + background + first | 90.428 |
| preference ~ task + background + active | 90.411 |
| preference ~ task + first + active | 89.657 |